

Introduction to Machine Learning

Homework 6: Support Vector Machines

Jack Langerman

1. Consider the data set for four points with features $\mathbf{x}_i = (x_{i1}, x_{i2})$ and binary class labels $y_i = \pm 1$.

x_{i1}	0	1	1	2
x_{i2}	0	0.3	0.7	1
y_i	-1	-1	1	1

- (a) Find a linear classifier that separates the two classes. Your classifier should be of the form

$$\hat{y} = \begin{cases} 1 & \text{if } b + w_1x_1 + w_2x_2 > 0 \\ -1 & \text{if } b + w_1x_1 + w_2x_2 < 0 \end{cases}$$

State the intercept b and weights w_1 and w_2 for your classifier. Note there is no unique answer as there are multiple linear classifiers that could separate the classes.

line at $x_{i2} = 0.5$ for all x_{i1} :

$$b = -0.5, w_1 = 0, w_2 = 1$$

$$\hat{y} = \begin{cases} 1 & \text{if } b + w_1x_1 + w_2x_2 > 0 \\ -1 & \text{if } b + w_1x_1 + w_2x_2 < 0 \end{cases}$$

- (b) Find the maximum γ such that $y_i(b + w_1x_{i1} + w_2x_{i2}) \geq \gamma$, for all i , for the classifier in part (a)?

$$\gamma = 0.5$$

- (c) Compute the margin of the classifier $m = \frac{\gamma}{\|\mathbf{w}\|}$, $\|\mathbf{w}\| = \sqrt{w_1^2 + w_2^2}$.

$$\|\mathbf{w}\| = \sqrt{w_1^2 + w_2^2} = \sqrt{0^2 + 1^2} = 1$$

$$m = \frac{\gamma}{\|\mathbf{w}\|} = \frac{0.5}{1} = 0.5$$

- (d) Which samples i are on the margin for your classifier?

Assuming $i = 0.3$ then samples $x_i, i = \{1, 2\}$ are in the margin.

$$(1)(0) + (0.3)(1) - 0.5 = -0.2$$

$$|-0.2| < 0.5$$

2. Consider the data set with scalar features x_i and binary class labels $y_i = \pm 1$.

x_i	0	1.3	2.1	2.8	4.2	5.7
y_i	-1	-1	-1	1	-1	1

Consider a linear classifier for this data of the form,

$$\hat{y} = \begin{cases} 1 & z > 0 \\ -1 & z < 0, \end{cases} \quad z = x - t,$$

where t is a threshold. For each threshold t , let $J(t)$ denote the sum hinge loss,

$$J(t) = \sum_i \epsilon_i, \quad \epsilon_i = \max(0, 1 - y_i z_i).$$

- (a) Write a short python program to plot $J(t)$ vs. t for 100 values of t in the interval $t \in [0, 5]$.

```
x = np.array([0, 1.3, 2.1, 2.8, 4.2, 5.7])
y = np.array([-1, -1, -1, 1, -1, 1])

def predict(x, t):
    z = x-t
    # 1 if z>0, -1 if z<0, break if z==0
    return 1 if z > 0 else (-1 if z < 0 else float('NaN'))

def hinggeLoss(x, y, t):
    z = x-t
    epsilon = np.maximum(0, 1 - y*z)
    return np.sum(epsilon)

plt.figure(figsize=(12,4))
for i, cmp in enumerate([lambda a, b: a<b, lambda a, b: a<=b]):
    plt.subplot(1,2,i+1)
    hist = {}
    minLoss = float('Inf')
    t_best = -1
    hist['t'] = []
    hist['J'] = []

    for t in np.linspace(0, 5, 100):
        J = hinggeLoss(x, y, t)

        if cmp(J, minLoss):
            minLoss = J
            t_best = t

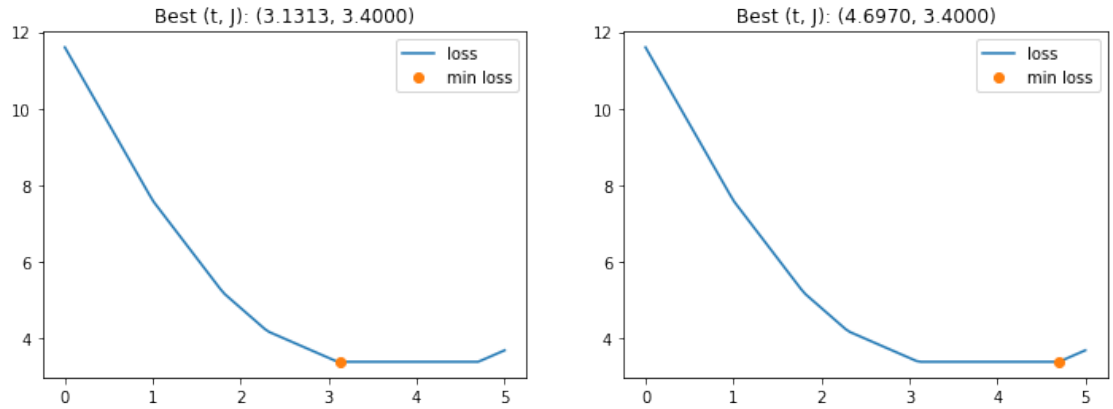
    hist['t'].append(t)
```

```

hist['J'].append(J)

plt.title("Best (t, J): ({:.4f}, {:.4f})".format(t_best, minLoss))
plt.plot(hist['t'], hist['J'], label="loss")
plt.plot(t_best, minLoss, 'o', label="min loss")
plt.legend()

```



(b) Based on the plot, what is one value of t that minimizes $J(t)$.

$$t = 4.5$$

(c) For the value of t in part (b), find the corresponding slack variables ϵ_i .

```

t = 4.5
z = x-t
epsilon = np.maximum(0, 1 - y*z)

print(epsilon)

```

```
[ 0  0  0  2.7  0.7  0 ]
```

(d) Which samples i violate the margin ($\epsilon_i > 0$) and which samples i are misclassified ($\epsilon_i > 1$).

x	y	y_hat	correct	slack
0.0	-1	-1	True	0.0
1.3	-1	-1	True	0.0
2.1	-1	-1	True	0.0
2.8	1	-1	False	2.7
4.2	-1	-1	True	0.7
5.7	1	1	True	0.0

3. Consider an image recognition problem, where an image \mathbf{X} and filter \mathbf{W} are 4×4 matrices:

$$\mathbf{X} = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}, \quad \mathbf{W} = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}.$$

- (a) Recall that in linear classification, the 4×4 image matrices \mathbf{X} and \mathbf{W} can be represented as 16-dimensional vectors, $\mathbf{x} = \text{vec}(\mathbf{X})$ and $\mathbf{w} = \text{vec}(\mathbf{W})$ by stacking the columns of the matrices vertically. What are \mathbf{x} and \mathbf{w} for the matrices above.

$$\mathbf{x} = [0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0]^T$$

$$\mathbf{w} = [0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0]^T$$

- (b) What is the inner product $z = \mathbf{w}^T \mathbf{x}$.

$$\mathbf{w}^T \mathbf{x} = 2$$

- (c) What is the inner product $z = \mathbf{w}^T \mathbf{x}_{\text{right}}$ where $\mathbf{x}_{\text{right}}$ is the vector corresponding to the matrix \mathbf{X} right shifted by one pixel with the left column filled with zeros.

$$\mathbf{w}^T \mathbf{x}_{\text{right}} = 0$$

- (d) What is the inner product $z = \mathbf{w}^T \mathbf{x}_{\text{left}}$ where \mathbf{x}_{left} is the vector corresponding to the matrix \mathbf{X} left shifted by one pixel with the right column filled with zeros.

$$\mathbf{w}^T \mathbf{x}_{\text{left}} = 2$$

- (e) Write the python command that can convert a 4×4 image matrix, \mathbf{X}_{mat} to the 16-dimensional vector, \mathbf{x} . What is the python command to go from \mathbf{x} to \mathbf{X}_{mat} .

```
def vec(A):
    B = []
    [B.extend(A[:, i]) for i in range(A.shape[1])]
    return np.array(B)
```

```
X = [
    [0, 0, 0, 0],
    [0, 0, 1, 0],
    [0, 0, 1, 0],
    [0, 0, 1, 0]
]
```

```
x = vec(X)
print("x =", x)
```

```
x = [0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0]
```

```
print(x.reshape(4,4, order='F'))
```

```
[[0 0 0 0]
 [0 0 1 0]
 [0 0 1 0]
 [0 0 1 0]]
```

4. Consider the data set with scalar features x_i and binary class labels $y_i = \pm 1$.


x_i	0	1	2	3
y_i	1	-1	1	-1

A support vector classifier is of the form


$$\hat{y} = \begin{cases} 1 & z > 0 \\ -1 & z < 0, \end{cases} \quad z = \sum_i \alpha_i y_i K(x_i, x),$$

where $K(x, x')$ is the radial basis function, $K(x, x') = e^{-\gamma(x-x')^2}$, and $\gamma > 0$ and $\alpha = [\alpha_1, \dots, \alpha_4]$ are parameters of the classifier.

- (a) Use python to plot z vs. x and \hat{y} vs. x when $\gamma = 3$ and $\alpha = [0, 0, 1, 1]$.



- (b) Repeat (a) with $\gamma = 0.3$ and $\alpha = [1, 1, 1, 1]$.



- (c) Which classifier makes more errors on the training data.

The classifier in part (b) makes more errors on the training set

HW 6 notebook

November 19, 2017

```
In [19]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt

%matplotlib inline

In [20]: x = np.array([0, 1.3, 2.1, 2.8, 4.2, 5.7])
y = np.array([-1, -1, -1, 1, -1, 1])

In [84]: def predict(x, t):
z = x-t
return 1 if z > 0 else (-1 if z < 0 else float('NaN')) # 1 if z>0, -1 if z<0, break
return res

In [85]: def hingedLoss(x, y, t):
z = x-t
epsilon = np.maximum(0, 1 - y*z)
return np.sum(epsilon)
```

0.0.1 Part 2 (a)

```
In [86]: plt.figure(figsize=(12,4))
for i, cmp in enumerate([lambda a, b: a<b, lambda a, b: a<=b]):
    plt.subplot(1,2,i+1)
    hist = {}
    minLoss = float('Inf')
    t_best = -1
    hist['t'] = []
    hist['J'] = []


    for t in np.linspace(0, 5, 100):
        J = hingedLoss(x, y, t)

        if cmp(J, minLoss):
            minLoss = J
            t_best = t

    hist['t'].append(t)
```

```
hist['J'].append(J)
```

```
plt.title("Best (t, J): ({:.4f}, {:.4f})".format(t_best, minLoss))
plt.plot(hist['t'], hist['J'], label="loss")
plt.plot(t_best, minLoss, 'o', label="min loss")
plt.legend()
```



0.0.2 Part 2 (c)

```
In [111]: t = 4.5
          z = x-t
          epsilon = np.maximum(0, 1 - y*z)
```

```
print(epsilon)
```

```
[ 0.  0.  0.  2.7  0.7  0.]
```

0.0.3 Part 2 (d)


```
In [112]: def predict(x, t):
          z = x-t
          res = np.zeros_like(z)
          res[np.greater(z, 0)] = 1
          res[np.less(z, 0)] = -1
          res[np.equal(z, 0)] = float('NaN')
          return res
```

```
In [120]: plt.plot(x, y, 'o')
          plt.plot([t,t], [-1, 1])
```

```
print(" x\ty\ty_hat\tcorrect\tslack")
```

```
print("", *["{\t}\t{\t}\t{\t}\t{\t}\n".format(*u) for u in zip(x, y, predict(x, t), predict(x, t), slack)])
```

x	y	y_hat	correct	slack
0.0	-1	-1.0	True	0.0
1.3	-1	-1.0	True	0.0
2.1	-1	-1.0	True	0.0
2.8	1	-1.0	False	2.7
4.2	-1	-1.0	True	0.7000000000000002
5.7	1	1.0	True	0.0



```
In [122]: print(" x & y & y_hat & correct & slack \\hline")
          print("", *["{} & {} & {} & {} & {} \\\\n".format(*u) for u in zip(x, y, predict(x, t), correct, slack)])

x & y & y_hat & correct & slack \\hline
0.0 & -1 & -1.0 & True & 0.0 \\\
1.3 & -1 & -1.0 & True & 0.0 \\\
2.1 & -1 & -1.0 & True & 0.0 \\\
2.8 & 1 & -1.0 & False & 2.7 \\\
4.2 & -1 & -1.0 & True & 0.7000000000000002 \\\
5.7 & 1 & 1.0 & True & 0.0 \\\
```

0.0.4 Part 3 (a)

```
In [158]: X = [
          [0, 0, 0, 0],
```



```

    [0, 0, 1, 0],
    [0, 0, 1, 0],
    [0, 0, 1, 0]
]

W = [
    [0, 0, 0, 0],
    [0, 1, 1, 0],
    [0, 1, 1, 0],
    [0, 0, 0, 0]
]

X, W = np.array(X), np.array(W)

def vec(A):
    B = []
    [B.extend(A[:, i]) for i in range(A.shape[1])]
    return np.array(B)

x = vec(X)
w = vec(W)

print("X =",x)
print("W =",w)
print("x \dot w =", np.dot(x, w))

X = [0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0]
W = [0 0 0 0 0 1 1 0 0 1 1 0 0 0 0 0]
x \dot w = 2

```

0.0.5 Part 3 (e)

```
In [159]: print(x.reshape(4,4, order='F'))
```

```

[[0 0 0 0]
 [0 0 1 0]
 [0 0 1 0]
 [0 0 1 0]]

```

0.0.6 Part 4 (a)

```

In [275]: # data
          x = np.array([0, 1, 2, 3])
          y = np.array([1, -1, 1, -1])

          # params

```

```

alpha = np.array([0, 0, 1, 1]) # dual vector
gamma = 3                      # param for rbf

def rbf(a, b):
    return np.exp(-gamma * (a-b)**2)

def predict(x, alpha, gamma, K):
    # data
    data_x = np.array([0, 1, 2, 3])
    data_y = np.array([1, -1, 1, -1])

    # score
    z = np.matmul(K(x, data_x), alpha*data_y)

    # prediction
    y_hat = np.zeros_like(z)
    y_hat[np.greater(z, 0)] = 1
    y_hat[np.less(z, 0)] = -1

    return y_hat

```

```

In [276]: # data
data_x = np.array([0, 1, 2, 3])
data_y = np.array([1, -1, 1, -1])

# params
alpha = np.array([0, 0, 1, 1]) # dual vector
gamma = 3                      # param for rbf

# kernal function
def K(a, b):
    return np.exp(-gamma * (a-b)**2)

# points to classify
x = np.linspace(0,3)[: , None]

# interior score
z = np.matmul(K(x, data_x), alpha*data_y)

# prediction
y_hat = predict(data_x[:,None], alpha, gamma, K)


# accuracy
acc_a = np.mean(np.equal(y_hat, data_y))
print("\n\t{}% accuracy\n".format(100*acc_a))

# plot
plt.plot(x, z)

```

75.0% accuracy

Out[276]: [<matplotlib.lines.Line2D at 0x109eb5e48>]



0.0.7 Part 4 (b)

```
In [277]: # data
data_x = np.array([0, 1, 2, 3])
data_y = np.array([1, -1, 1, -1])

# params
alpha = np.array([1, 1, 1, 1]) # dual vector
gamma = 0.3                    # param for rbf

# kernal function
def K(a, b):
    return np.exp(-gamma * (a-b)**2)

# points to classify
x = np.linspace(0,3)[: , None]

# interior score
z = np.matmul(K(x, data_x), alpha*data_y)
```

```

# prediction
y_hat = predict(data_x[:,None], alpha, gamma, K)


# accuracy
acc_b = np.mean(np.equal(y_hat, data_y))
print("\n\t{}% accuracy\n".format(100*acc_b))

# plot
plt.plot(x, z)

```

50.0% accuracy

Out[277]: [



0.0.8 Part 4 (c)

```

In [281]: print("\n\ta) {}% accuracy".format(100*acc_a))
          print("\tb) {}% accuracy\n".format(100*acc_b))
          print("\tThe settings in part (a) yield higher accuracy")

```

a) 75.0% accuracy

b) 50.0% accuracy

The settings in part (a) yield higher accuracy

In []: